

50 Likely Interview Questions & Answers — KPMG Forensic Data Analytics

Based on Nitish Pardal's resume + the KPMG JD

Answers are written in first person so you can adapt and practice them directly. Fill in real numbers/specifics where marked *[specify]* — vague answers are the #1 thing that costs candidates credibility in follow-ups.

A. Opening / Behavioral (Q1-8)

1. Tell me about yourself.

"I'm a Data Analyst with 3 years of experience across financial services and automotive manufacturing at Wipro. I work primarily in SQL, Python, and Power BI/Tableau — building dashboards, optimizing queries, and supporting fraud/risk monitoring. In my current role I've built fraud detection dashboards that contributed to a 15% reduction in fraud cases, and earlier in financial services I worked on ETL validation and anomaly detection for regulatory reporting. I'm looking to move into a Forensic analytics role because it lets me combine my SQL/BI skill set with the fraud and risk-monitoring work I've already been doing."

2. Why do you want to work at KPMG, specifically in Forensic Services?

"I've already been doing adjacent work — fraud dashboards, anomaly detection, risk monitoring — but at Wipro it was one part of a broader analyst role. KPMG's Forensic practice is where that becomes the core focus, with far more scale and structure — investigations, compliance, industry breadth. I want to go deeper into that specialization rather than treat it as one dashboard among many."

3. Walk me through your resume.

Walk chronologically: B.Tech CS/IT (2018-22) → Junior Analyst in Financial Services (Mar'23-Nov'23): ETL validation, DB2 optimization, fraud/risk monitoring → Data Analyst in Automotive Manufacturing (Dec'23-present): SQL optimization, Power BI/Tableau dashboards, fraud detection visualizations, predictive analytics. Keep it under 90 seconds — this is a warm-up question, not the main event.

4. Why are you looking to leave your current role?

Keep this positive, forward-facing, never negative about Wipro: "I've grown a lot in my current role, but I've hit the point where I want to specialize — go deeper into fraud/forensic analytics and cloud-based data engineering rather than general BI reporting. KPMG's Forensic practice is exactly that specialization."

5. What are your strengths?

Anchor in resume evidence: "SQL optimization is my strongest technical skill — I've delivered 30-40% performance improvements in two different roles. I'd also point to translating messy business problems into dashboards non-technical stakeholders can act on — that's what the fraud dashboard project was really about."

6. What's a weakness you're working on?

Pick something real and JD-relevant, with an active fix: "My hands-on cloud experience is narrower than I'd like — I've worked with BigQuery/GCP but haven't owned pipeline infrastructure or IaC tools like Terraform. I'm actively studying this because I know it's core to this role."

7. Where do you see yourself in 3-5 years?

"Growing into a specialist forensic data analyst — eventually leading investigations or analytics workstreams end-to-end, and building deeper expertise in fraud analytics techniques and cloud data engineering."

8. Why should we hire you over a candidate with more years of experience?

"I'd point to the density of what I've done in 3 years — I've already delivered the exact things this JD asks for at a smaller scale: fraud dashboards with measurable impact, SQL optimization with measurable impact, cross-functional work with risk and finance teams. I'm looking for the room to do that at KPMG's scale, and I'll be very deliberate about closing the cloud/DevOps gap fast."

B. Experience Gap (Q9-10)

9. The JD asks for 4-7 years — you have 3. Why did you apply?

"I was upfront about this with myself before applying — I have 3 years, not 4-7. What I'd point to is that I've already been doing Assistant Manager-level work: end-to-end ownership of dashboards that changed business outcomes, not just executing tickets. I think I can close the gap in scope faster than in years, and I'd rather be evaluated on what I've delivered than just the number on the calendar."

10. Do you feel ready for a role that expects independent, end-to-end ownership of business cases?

"Largely yes — my fraud dashboard work and ETL validation work were both things I owned start to finish, from talking to risk/finance stakeholders to shipping the dashboard. What I haven't owned yet is full pipeline/cloud infrastructure — that's the part I'd be leaning on the team to ramp up on quickly."

C. SQL — Resume & Technical Depth (Q11-20)

11. Walk me through your SQL query optimization work (30% and 40% improvements).

"At my current role I improved data retrieval by 30% mainly through restructuring joins and reducing unnecessary subqueries. At my previous role, DB2 performance improved 40% through indexing strategy and query restructuring — for example replacing correlated subqueries with joins and adding indexes on frequently filtered columns."

12. What's the difference between WHERE and HAVING?

"WHERE filters rows before aggregation; HAVING filters groups after aggregation. You can't use an aggregate function like COUNT() or AVG() in a WHERE clause — that has to go in HAVING."

13. Correlated subquery vs JOIN — which is faster and why?

"A correlated subquery re-executes once per outer row, which can be very slow on large tables. A JOIN lets the optimizer process the whole set at once, usually via a hash or merge join, which scales much better. I'd rewrite correlated subqueries into joins wherever possible — that's actually one of the optimization techniques I used in my DB2 work."

14. How do you find duplicate records in a table?

"I use ROW_NUMBER() partitioned by the columns that define a duplicate, ordered by some tiebreaker like a timestamp or ID, then filter WHERE row_num > 1." *(Be ready to write this out — see your practice queries from earlier.)*

15. Difference between ROW_NUMBER(), RANK(), and DENSE_RANK()?

"ROW_NUMBER() always assigns a unique number even to ties, which is why it's the right tool for isolating duplicate rows. RANK() gives tied rows the same rank but skips the next number. DENSE_RANK() also ties but doesn't skip — the next rank is always sequential."

16. How would you delete duplicate rows from a table safely?

"First I'd SELECT with ROW_NUMBER() to see exactly which rows would be affected — never delete blind. Then in SQL Server I'd wrap that in a CTE and DELETE FROM the CTE where row_num > 1, ideally inside a transaction so I can roll back if needed, and verify row counts before and after."

17. What's the difference between CREATE TABLE AS SELECT, SELECT INTO, and INSERT INTO SELECT?

"CREATE TABLE AS SELECT and SELECT INTO both create a new table and populate it — SELECT INTO is the SQL Server-specific syntax. INSERT INTO SELECT requires the target table to already exist — it only populates, it doesn't create."

18. How do you approach optimizing a slow-running query?

"I start with the execution plan to see where the time is going — usually table scans instead of index seeks. Then I look at whether indexes exist on filter/join columns,

whether SELECT * is pulling unnecessary columns, and whether subqueries can be rewritten as joins. I did exactly this to get the DB2 improvement in my last role."

19. Write a query to flag transactions more than 2 standard deviations from a customer's average. (forensic-flavored — practice this one specifically)

```
sql

WITH stats AS (
    SELECT CustomerID, AVG(Amount) AS avg_amt, STDEV(Amount) AS std_amt
    FROM Transactions
    GROUP BY CustomerID
)
SELECT t.*
FROM Transactions t
JOIN stats s ON t.CustomerID = s.CustomerID
WHERE ABS(t.Amount - s.avg_amt) > 2 * s.std_amt;
```

20. What databases have you worked with, and how do they differ in your experience?

"MySQL and BigQuery in my current role, DB2 in financial services. DB2 required more manual indexing discipline for performance; BigQuery is columnar and optimized differently — for large scans rather than transactional lookups, so I approach query design differently there (avoiding SELECT *, being mindful of partitioning)."

D. Python (Q21-26)

21. How do you use Python alongside SQL in your work?

"SQL handles the heavy extraction/aggregation on the database side; I use Python — mainly Pandas and NumPy — for cleaning, transforming, and shaping the data further once it's out of the DB, especially when the transformation logic is too complex or iterative for SQL alone."

22. How would you detect outliers in a dataset using Python?

"With Pandas, I'd typically calculate z-scores or use IQR (Q1/Q3 and 1.5×IQR bounds) to flag values outside the normal range. For something like fraud analytics I'd combine that with grouping — e.g., outliers relative to a customer's own historical average, not just the global average."

23. Merge vs Join in Pandas — what's the difference?

"They're functionally very similar — `pd.merge()` is the general function and `.join()` is a DataFrame method optimized for joining on index. I default to `merge()` since it's more explicit about join type (inner/left/right/outer) and join keys."

24. How do you handle missing data in Pandas?

"Depends on context — I'd check the proportion missing first. For small gaps I might use `fillna()` with mean/median or forward-fill for time series; for large or systematic gaps I'd flag and investigate rather than blindly imputing, since that can mask a real data quality issue — which matters a lot in a forensic context."

25. Tell me about your Customer Spending Analysis Python project.

Reference resume: "I cleaned, processed, and visualized 100K+ transactions using Pandas and Seaborn to identify purchasing trends." Be ready to describe: what cleaning steps (nulls, duplicates, type conversions), what visualizations (distribution plots, trend lines), and one specific insight it surfaced.

26. Have you used Python for automation? How?

Tie to resume: "I automated recurring reporting workflows, which cut manual reporting effort by 40%. That involved scripting the data pull/transform steps that used to be done manually, then feeding cleaned output into the Power BI/Tableau refresh."

E. Power BI / Tableau / Dashboards (Q27-33)

27. Walk me through the fraud detection dashboard you built.

This is your headline story — see the full STAR breakdown in your prep doc from earlier. Situation → risk/ops needed faster fraud visibility → built Power BI dashboard with DAX-based flags on transaction/production anomalies → 15% reduction in fraud cases. Be ready for the two guaranteed follow-ups below (Q28, Q29).

28. What specific rules or patterns did the fraud dashboard flag?

Have 2-3 concrete examples ready: unusual transaction volume/frequency, deviation from historical average (tie to Q19's stdev logic), duplicate or mismatched entries across systems. If your actual dashboard was simpler than this, be honest and add what you'd enhance it with if asked to make it more forensic-grade.

29. How did you validate flagged cases weren't false positives?

"I worked with risk/finance teams to review flagged cases and tightened thresholds over time based on their feedback — a purely statistical flag isn't enough on its own without business context."

30. Calculated column vs measure in Power BI — what's the difference?

"A calculated column is computed row-by-row and stored in the table, using row context. A measure is computed on the fly based on the current filter context in a visual, and isn't stored — I default to measures wherever possible since they're more memory-efficient and responsive to slicers/filters."

31. What DAX functions do you use most often?

Mention from your actual work: CALCULATE, FILTER, SUMX, time intelligence functions if used. Be ready with one real example, e.g., a rolling count/average over a time window (similar to the Medium article's `RechargeCount1hr` pattern) — good to reference if it fits your dashboard's logic.

32. How do you decide what to visualize for a non-technical stakeholder?

"I start with what decision they need to make, not what data is available. For the fraud dashboard, risk/ops needed to quickly spot 'is this unusual' — so KPI cards and trend lines with clear flagging mattered more than dense tables."

33. Power BI vs Tableau — differences in your experience?

"Power BI integrates more tightly with the Microsoft stack and DAX gives strong control over business logic. Tableau's strength is more intuitive drag-and-drop visual exploration, particularly for ad hoc analysis. I've used both, and choose based on the audience and existing tooling in the org."

F. Fraud / Forensic Specific (Q34-40)

34. What's your understanding of KPMG's Forensic practice?

"It covers both reactive work — investigating fraud, misconduct, tracing funds and assets — and preventive work like building monitoring systems, using accounting, technology, and industry expertise across sectors like finance, healthcare, and media."

35. Have you worked on fraud investigations, or just fraud-adjacent monitoring/reporting?

Be honest: your resume shows monitoring/detection dashboards, not formal investigations. "My experience has been building the detection/monitoring layer — dashboards and anomaly flags — rather than running formal investigations. I see this role as the natural next step into that deeper investigative side."

36. What statistical techniques are relevant to fraud detection that you're aware of?

Mention (even if not all hands-on): z-score/IQR outlier detection, Benford's Law for digit-pattern anomalies in financial figures, clustering to group similar anomalous records, rule-based thresholds. Be honest about which you've actually applied vs know conceptually.

37. How would you approach investigating a sudden spike in a business metric (e.g., claims, transactions)?

"First I'd segment — is the spike broad or concentrated in one region/customer/product? Then I'd check for a legitimate driver (seasonality, a promotion) before assuming anomaly. If unexplained, I'd drill into the underlying records for patterns — timing, duplicate entities, unusual amounts — the same root-cause approach I used for data discrepancies in financial reporting."

38. What red flags would indicate financial statement fraud?

General knowledge answer: unusual revenue recognition timing, round-number transactions, transactions just under approval thresholds, unexplained related-party transactions, inconsistent trends vs industry peers. Frame as "things I understand conceptually and would want to develop hands-on judgment on with this team."

39. Have you worked with any AML (anti-money laundering) or compliance-related datasets?

If not directly, connect to closest experience: "Not AML specifically, but my financial services work involved regulatory reporting and anomaly detection on transaction data, which is conceptually adjacent."

40. Why does a "good to have" like Generative AI exposure matter for forensic analytics?

"GenAI can help triage large volumes of unstructured data — documents, communications — much faster than manual review, which matters a lot in investigations. I haven't used it in a forensic context yet, but I'm comfortable with the underlying analytics mindset and keen to pick it up."

G. Cloud / Data Engineering — JD Gap Area (Q41-45)

41. What cloud platforms have you worked with?

"Mainly GCP/BigQuery — querying and analyzing data rather than building infrastructure. I understand data lake vs data warehouse concepts, but I haven't personally provisioned cloud resources or built pipelines end-to-end."

42. Data lake vs data warehouse — what's the difference?

"A data lake stores raw, often unstructured or semi-structured data cheaply at scale, schema-on-read. A data warehouse stores structured, cleaned, modeled data optimized for fast querying and reporting, schema-on-write."

43. Have you used Terraform or CloudFormation?

Honest, no bluffing: "No hands-on experience there. I understand the concept — defining infrastructure as code instead of manual provisioning — but haven't used either tool directly. It's an area I'd want to ramp up quickly in this role."

44. What is CI/CD, and have you worked with any pipelines?

"CI/CD automates testing and deployment so changes ship reliably and frequently. I've used Git/GitHub for version control, but I haven't set up or managed CI/CD pipelines myself — that's more on the engineering side of what I've done so far."

45. How would you ensure data security and governance in a cloud environment?

General-knowledge answer, tie to what you do know: "Principles I understand — access controls/least privilege, encryption at rest and in transit, audit logging, and clear data classification policies. My hands-on governance experience has been more on the data-quality/compliance side — validation, reconciliation, anomaly detection — than infrastructure-level security."

H. Agile / Process (Q46-47)

46. Have you worked in an Agile/Scrum environment?

"I've worked in JIRA-tracked environments with sprint-based work, though I wasn't in a formal Scrum role. I'm familiar with the ceremonies — stand-ups, sprint planning, retrospectives — from participating in them, even if I haven't led them."

47. How do you handle competing priorities or tight deadlines across projects?

Use a real example if you have one — e.g., balancing dashboard delivery with ongoing SQL optimization work. Structure: what you did to triage, how you communicated with stakeholders, the outcome.

I. Closing / Culture Fit (Q48-50)

48. How do you handle disagreement with a stakeholder about your analysis or findings?

"I'd walk them through the data and methodology first — often disagreement comes from a gap in shared understanding, not a real conflict. If they still see it differently after that, I'd look for what additional data or context could resolve it rather than just deferring or digging in."

49. What questions do you have for us?

Use the 2-3 from your prep doc: ask about reactive vs preventive work split, the actual cloud tech stack day-to-day, and what success looks like in the first 6 months for someone at your experience level.

50. What are your salary expectations?

Don't box yourself in with a hard number this early unless pressed: "I'd like to understand the full scope of the role and comp structure first, but I'm open to discussing a number aligned with market standard for this level — happy to share a range if that's helpful." *(If pressed for a number, have a researched range ready beforehand — check current Consultant/AM data analytics salary bands in Bangalore/Delhi NCR.)*

- Don't try to memorize all 50 word-for-word — that reads as robotic.
- Prioritize: **Section E (Fraud/Forensic)** and **Section C/D (SQL/Dashboard depth)** first — these are where you'll be judged hardest given the role.
- For Section F (Cloud/DevOps gap), practice saying the honest answers *confidently* — a shaky "I don't know" reads worse than a calm "not yet, and here's my plan to close that."
- Say a few of these out loud, not just read them — tone matters as much as content tomorrow.